

AFISHA AZME

Data Scientist • Machine Learning Engineer • Python | SQL | Power BI

+91-6303662944 • azmeshaik@gmail.com • [linkedin.com/in/afisha-azme-612400357](https://www.linkedin.com/in/afisha-azme-612400357) • github.com/afisha-02 •
Bengaluru / Hyderabad

PROFESSIONAL SUMMARY

Data Scientist with 6 months of hands-on experience building end-to-end ML pipelines and analyzing datasets up to 110K+ records. Improved minority-class F1-score by 18% and delivered a segmentation-driven strategy projected to reduce patient no-show rates by 15–20%. Skilled in Python, SQL, Machine Learning, Power BI, and NLP/GenAI with a strong focus on data analysis, business insights, and stakeholder-ready reporting. Seeking full-time Data Scientist / ML Engineer roles in Bengaluru or Hyderabad. Available for immediate joining.

CORE COMPETENCIES

Data Analysis & EDA	ML Model Development	SQL & Data Extraction	Business Insights	Data Cleaning & ETL	NLP & GenAI
Python (Pandas, NumPy)	XGBoost / Random Forest	Power BI & Tableau	Stakeholder Reporting	MLflow & MLOps	Feature Engineering

TECHNICAL SKILLS

Languages & Querying: Python (Pandas, NumPy, scikit-learn), SQL

Machine Learning: Classification, Clustering, Regression, Feature Engineering, XGBoost, Random Forest, Model Evaluation (AUC-ROC, F1, Precision/Recall)

NLP & GenAI: TF-IDF, Sentence Embeddings, LangChain, LLM Integration, Prompt Engineering

Data Visualization: Matplotlib, Seaborn, Plotly, Power BI, Tableau

MLOps & Tools: MLflow, Jupyter Notebook, Git, Google Colab, Excel, SMOTE

Statistics: Hypothesis Testing, Descriptive & Inferential Statistics, EDA

PROFESSIONAL EXPERIENCE

Data Science Intern | **Edukron Technologies** • Remote Oct 2025 – Mar 2026

- Built ML classification pipelines using Random Forest and XGBoost across 5+ real-world datasets; applied SMOTE oversampling and 5-fold cross-validation, improving minority-class F1-scores by up to 18% on imbalanced data — directly reducing model error on high-priority business cases.
- Designed end-to-end NLP pipelines (tokenization, TF-IDF, sentence embeddings, sentiment & topic classification with LangChain), enabling automated text analysis at scale and reducing manual review effort.
- Developed 3+ interactive dashboards (Plotly, Seaborn) visualizing model performance and EDA findings, measurably accelerating stakeholder decision-making speed across internal projects.
- Executed comprehensive EDA on structured datasets of 10K–100K rows — including missing-value imputation, outlier detection, and feature selection — consistently improving downstream model accuracy.
- Maintained structured MLflow experiment logs and reproducible documentation for all model iterations, enabling independent result replication and reducing team onboarding friction.

PROJECTS

Healthcare Appointment No-Show Prediction — Python | scikit-learn | XGBoost | Pandas | Seaborn

- Cleaned and analyzed 110,000+ patient appointment records; identified 5 key predictors of no-shows (SMS receipt, lead time, neighborhood, scholarship status, chronic conditions) through rigorous EDA.
- Trained and benchmarked Logistic Regression, Random Forest, and XGBoost classifiers; optimized recall to 78% to minimize missed high-risk patients, directly supporting operational scheduling decisions.
- Delivered a segmentation-based reminder strategy projected to reduce no-show rates by 15–20% for high-risk cohorts — translating model output into actionable business recommendations for hospital operations.

Customer Segmentation — RFM Clustering — *Python | K-Means | Pandas | Matplotlib*

- Applied K-Means clustering ($k=4$, Silhouette Score = 0.61) on RFM features derived from 8,000+ customer transactions, identifying 4 distinct behavioral segments: Champions, At-Risk, Hibernating, and New Customers.
- Formulated tailored CRM and retention strategies per segment to improve marketing ROI, bridging analytical findings with business strategy for non-technical stakeholders.
- Presented cluster distributions and feature importance in a stakeholder-ready slide deck, communicating data insights clearly to non-technical audiences.

Insurance Claim Cost Prediction — *Python | scikit-learn | XGBoost | Random Forest | Neural Network | Kaggle Dataset*

- Built a supervised regression pipeline to predict individual insurance claim costs from patient health data; performed thorough EDA and feature engineering to surface strongest cost drivers.
- Benchmarked 5 models (Linear Regression, Decision Tree, Random Forest, XGBoost, Neural Network) on R^2 and RMSE; selected best model via cross-validation and hyperparameter tuning.
- Identified smoking status and BMI as dominant predictors of high claim costs; visualized feature importance and cost distribution patterns to communicate findings clearly to academic reviewers.

EDUCATION

B.Tech — Computer Science & Engineering | JNTUA, Anantapur *Apr 2026 • CGPA: 7.5*

CERTIFICATIONS

- Oracle Cloud Infrastructure 2025 Certified Generative AI Professional — Oracle
- Introduction to Prompt Engineering with GitHub Copilot — GitHub / Microsoft
- Gemini for Google Workspace — Google
- Artificial Intelligence Fundamentals — IBM

Available for immediate joining | Open to relocation | Actively seeking opportunities